

# Intro to Data Science - HW 4

Copyright 2022, Jeffrey Stanton, Jeffrey Saltz, and Jasmina Tacheva

```
# Enter your name here: Bhavya Shah
```

Attribution statement: (choose only one and delete the rest)

```
# 1. I did this homework by myself, with help from the book and the professor.
```

## (Chapter 6 of Introduction to Data Science)

Reminders of things to practice from previous weeks:

Descriptive statistics: `mean()` `max()` `min()`

Sequence operator: `:` (For example, `1:4` is shorthand for `1, 2, 3, 4`)

Create a function: `myFunc <- function(myArg) { }`

?command: Ask R for help with a command

**This module: Sampling** is a process of **drawing elements from a larger set**. In data science, when analysts work with data, they often work with a sample of the data, rather than all of the data (which we call the **population**), because of the expense of obtaining all of the data.

One must be careful, however, because **statistics from a sample rarely match the characteristics of the population**. The **goal of this homework** is to **sample from a data set several times and explore the meaning of the results**. Before you get started make sure to read Chapter 6 of *An Introduction to Data Science*. Don't forget your comments!

## Part 1: Write a function to compute statistics for a vector of numeric values

- A. Create a new function which takes a numeric vector as its input argument and returns a dataframe of statistics about that vector as the output. As a start, the dataframe should have the **min**, **mean**, and **max** of the vector. The function should be called **statsCalc**:

```
statsCalc<-function(a){  
  b<-c(a)  
  df<-data.frame(b)  
  df$min<-min(df$b)  
  df$mean<-mean(df$b)  
  df$max<-max(df$b)  
  df$median<-median(df$b)  
  df$standard_deviation<-sd(df$b)  
  return(df)  
}
```

- B. Test your function by calling it with the numbers **one through ten**:

```
statsCalc(1:10)
```

```
##      b min mean max median standard_deviation
## 1    1  1  5.5  10    5.5          3.02765
## 2    2  1  5.5  10    5.5          3.02765
## 3    3  1  5.5  10    5.5          3.02765
## 4    4  1  5.5  10    5.5          3.02765
## 5    5  1  5.5  10    5.5          3.02765
## 6    6  1  5.5  10    5.5          3.02765
## 7    7  1  5.5  10    5.5          3.02765
## 8    8  1  5.5  10    5.5          3.02765
## 9    9  1  5.5  10    5.5          3.02765
## 10  10  1  5.5  10    5.5          3.02765
```

C. Enhance the statsCalc() function to add the **median** and **standard deviation** to the returned dataframe.

```
statsCalc(1:10)
```

```
##      b min mean max median standard_deviation
## 1    1  1  5.5  10    5.5          3.02765
## 2    2  1  5.5  10    5.5          3.02765
## 3    3  1  5.5  10    5.5          3.02765
## 4    4  1  5.5  10    5.5          3.02765
## 5    5  1  5.5  10    5.5          3.02765
## 6    6  1  5.5  10    5.5          3.02765
## 7    7  1  5.5  10    5.5          3.02765
## 8    8  1  5.5  10    5.5          3.02765
## 9    9  1  5.5  10    5.5          3.02765
## 10  10  1  5.5  10    5.5          3.02765
```

D. Retest your enhanced function by calling it with the numbers **one through ten**:

Note that the code below has an error, so just running the code will not work. Fix the code and then run test function.

```
statsCalc(1:10)
```

```
##      b min mean max median standard_deviation
## 1    1  1  5.5  10    5.5          3.02765
## 2    2  1  5.5  10    5.5          3.02765
## 3    3  1  5.5  10    5.5          3.02765
## 4    4  1  5.5  10    5.5          3.02765
## 5    5  1  5.5  10    5.5          3.02765
## 6    6  1  5.5  10    5.5          3.02765
## 7    7  1  5.5  10    5.5          3.02765
## 8    8  1  5.5  10    5.5          3.02765
## 9    9  1  5.5  10    5.5          3.02765
## 10  10  1  5.5  10    5.5          3.02765
```

## Part 2: Sample repeatedly from the New York State COVID Testing Dataset from HW 3

- A. Load the dataset from the following URL, using `read_csv`: [https://data-science-intro.s3.us-east-2.amazonaws.com/NYS\\_COVID\\_Testing.csv](https://data-science-intro.s3.us-east-2.amazonaws.com/NYS_COVID_Testing.csv) ([https://data-science-intro.s3.us-east-2.amazonaws.com/NYS\\_COVID\\_Testing.csv](https://data-science-intro.s3.us-east-2.amazonaws.com/NYS_COVID_Testing.csv))

```
library(tidyverse)
```

```
## — Attaching packages — tidyverse 1.3.2 —
## ✓ ggplot2 3.4.0      ✓ purrr   1.0.1
## ✓ tibble  3.1.8      ✓ dplyr   1.0.10
## ✓ tidyr   1.3.0      ✓ stringr 1.5.0
## ✓ readr   2.1.3      ✓ forcats 1.0.0
## — Conflicts — tidyverse_conflicts() —
## ✗ dplyr::filter() masks stats::filter()
## ✗ dplyr::lag()     masks stats::lag()
```

```
testDF<-data.frame(read.csv("https://data-science-intro.s3.us-east-2.amazonaws.com/NYS_COVID_Testing.csv"))
```

- B. Use `head(testDF)` and `tail(testDF)` to show the data. Add a comment that describes what each variable in the data set contains.

```
head(testDF)
```

```
##   TestDate AgeGroup PositiveCases TotalTests AgeCategory
## 1 3/2/2020 45 to 54             1         1 middle-aged_adults
## 2 3/3/2020 25 to 34             0         2   young_adults
## 3 3/3/2020 35 to 44             0         1 middle-aged_adults
## 4 3/3/2020 45 to 54             0         1 middle-aged_adults
## 5 3/3/2020 55 to 64             0         2   senior_citizens
## 6 3/3/2020 65 to 74             0         2   senior_citizens
```

```
tail(testDF)
```

```
##   TestDate AgeGroup PositiveCases TotalTests AgeCategory
## 7378 1/3/2022  5 to 19           9923     38977      children
## 7379 1/3/2022 55 to 64           5739     27019 senior_citizens
## 7380 1/3/2022 65 to 74           2759     14498 senior_citizens
## 7381 1/3/2022 75 to 84           1141       6519 senior_citizens
## 7382 1/3/2022   85 +             680       4028 senior_citizens
## 7383 1/3/2022    < 1             717       2074      children
```

```
#test date is the date when the test was done
#age group is different groups of age ranges
#positivecases is the number of positive cases on a particular date
#totalcases is the number of total test cases on a particular date
#agecategory is different age categories as p-er the age group
```

C. Sample ten observations from **testDF\$TotalTests**.

```
sampltotal<-sample(testDF$TotalTests,size=10,replace=TRUE)
```

D. Call your statsCalc( ) function with a new sample of ten observations from **testDF\$TotalTests**, where the sampling is done inside the **statsCalc** function call.

```
statsCalc(sampltotal)
```

```
##          b min   mean   max median standard_deviation
## 1    9314 396 9593.3 17951  10983          6203.833
## 2   13484 396 9593.3 17951  10983          6203.833
## 3     396 396 9593.3 17951  10983          6203.833
## 4   17951 396 9593.3 17951  10983          6203.833
## 5    1676 396 9593.3 17951  10983          6203.833
## 6    2651 396 9593.3 17951  10983          6203.833
## 7    8636 396 9593.3 17951  10983          6203.833
## 8   12652 396 9593.3 17951  10983          6203.833
## 9   12896 396 9593.3 17951  10983          6203.833
## 10 16277 396 9593.3 17951  10983          6203.833
```

E. Now use the **mean()** function, with another sample done inside the mean function. Is the mean returned from the **statsCalc** function the same as the mean returned from the mean function on this sample? Why or why not? Explain.

```
mean(sample(testDF$TotalTests,size=10,replace=TRUE))
```

```
## [1] 15204.8
```

```
#the mean we calculated by statscalc is different from the above mean as the sampling method is
random and the replacement value is true. Hence, 2 different means were obtained.
```

F. Use the **replicate( )** function to repeat your sampling of **testDF\$TotalTests** twenty times, with each sample calling **mean()** on ten observations. The first argument to **replicate( )** is the number of repeats you want. The second argument is the little chunk of code you want repeated.

```
replicate(20,mean(sample(testDF$TotalTests,size=10,replace=TRUE)),simplify = TRUE)
```

```
## [1] 9983.7 13942.3 13920.4 13127.0 16226.2 6175.2 21217.9 8741.8 9886.8
## [10] 14683.6 8106.8 7008.6 16996.5 12921.0 16823.6 15109.2 8278.2 13399.5
## [19] 7809.5 13796.9
```

G. Write a comment describing why every replication produces a different result.


*#sampling is process in which random elements are chosen from a population, and hence every draw n gives every element a fair chance and thus every time a different result is produced.*

H. Rerun your replication, this time doing 20 replications and storing the output of **replicate()** in a variable called **values**.

```
values<-replicate(20,mean(sample(testDF$TotalTests,size=10,replace=TRUE)),simplify = TRUE)
```

I. Generate a **histogram** of the means stored in **values**.


```
hist(values)
```



J. Repeat the replicated sampling, but this time, raise your replications to **1000**.

```
values1<-replicate(1000,mean(sample(testDF$TotalTests,size=10,replace=TRUE)),simplify = TRUE)  
hist(values1)
```

## Histogram of values1



K. Compare the two histograms - why are they different? Explain in a comment.

*#We have taken histogram for two samples, one with 20 replications and other with 1000. Since, while replication the values drawn can be random and hence different from each other. Thus, both the histograms are different from each other.*